

Java Packaging. Containers (1)

- Canonical Java Packaging
 - JAR Files. Fat JARs
 - JPMS + jlink
 - GraalVM native-image
- Linux Packaging (DEB, RPM)
- Virtualization
- Containers: Concepts, Tools, Implementation
- Container Orchestration
- Some Container Solutions

Deadline-Driven Development™

Apr 23: Second Release

- All previously found **defects fixed**
- **At least 2** User Stories implemented
- Proper **error handling**
- **Unit Tests** for all classes except framework-related
 - Line/Statement **Coverage $\geq 70\%$**

Canonical Java Packaging: JAR

- Just fancy **ZIP archives** containing **compiled Java classes**, **resources** and **metainformation** (META-INF/)
@see <https://docs.oracle.com/javase/7/docs/technotes/guides/jar/jar.html>
- **Compiled classes** and **class resources** are put into directories corresponding to Java packages.
E.g. class ru.hse.java.HelloWorld => ru/hse/java/HelloWorld.class
 - No directories are created for anonymous, inner and static inner classes, because they are synthesized by the compiler
E.g. class ru.hse.java.HelloWorld.MyCoolClass => ru/hse/HelloWorld\$MyCoolClass.class
- Most important **metainformation** is the Manifest, META-INF/MANIFEST.MF:
Manifest-Version: 1.0
Main-Class: ru.hse.MyCoolProjectMain
<more key-value pairs...>
—————▶ **java -jar my-cool-project-1.0.jar**
- META-INF/ directory also MAY contain:
 - Digital signature files (*.RSA, *.DSA, SIG-*)
 - **Service Provider** definitions (META-INF/services/<FQCN of Service Class Impl>). @see future seminar on DI

Fat (Uber) JAR

<https://stackoverflow.com/a/36539885/3438672>

- Bring all of your dependencies (transitively) in a single large JAR file
- Three methods:
 - **Unpack all dependency JARs** (`maven-assembly-plugin` → `jar-with-dependencies`)
 - Unpack all dependency JARs, but also rename their packages and merge resources (`maven-shade-plugin`)
 - To avoid conflicts with your library users' dependencies
 - Copy JARs into your JAR and use a special *class loader* to transparently access their classes (`onejar`, `spring-boot-plugin:repackage`)
- Classpath is Linear → JAR Hell
 - `maven-enforcer-plugin` helps, use it!

JPMS + jlink

<https://www.baeldung.com/java-9-modularity>

– Java Platform Module System

- Modules depend on each other (**requires** [transitive] ...)
- Modules have Strong Encapsulation™
 - Explicitly define which packages are visible to the outside world (**exports** ... [to ...])
 - Explicitly define which packages they allow reflection access to (**opens** ... [to ...])
- Dependency Cycles and Split Packages disallowed
- Inter-module dependencies are assumed to be *mostly* static (determined at app start time)
- Requires you to **cleanly separate your system into modules**, which is **NO SMALL TASK**
- But as a bonus, you can package only **modules really used** by your application:

```
jlink --module-path $JAVA_HOME/jmods:mllib \  
      --add-modules com.greetings \  
      --output greetingsapp
```

...the image will include **only used** JDK modules (**jmods:...**)

GraalVM native-image

- Ahead-of-Time compiles your Java code into a **native application for the target platform**
 - Some limitations, mostly around Reflection usage and static class initialization (**static finals** and **static {}** blocks spawning threads and the like)
 - Popular frameworks are supported out-of-box, less popular can encounter problems
- Emerging frameworks using native-image as the default: Quarkus, Micronaut
- Good fit for Microservices, Serverless, CLIs, Agents, ...
 - **Fast start** (low start latency) far more important than **peak throughput**
 - GC and runtime is naive compared to OpenJDK's Hotspot


Linux Packaging (RPM, DEB)

- General tools for managing components of a Linux distribution
 - Manage *packages*, which provide *files* installed into your filesystem
 - Packages are *versioned* and *depend* on each other (possibly with a version/version range specified)
 - Package Managers provide dependency resolution, download and installation
- Cannot have multiple versions of the same package
 - Multiple packages, however, can provide the same binary, e.g. `/bin/java` via `update-alternatives`

Virtualization

- Run 1+ *Guests* (OS w/virtualized hardware) on a single *Host*, managed by a *Hypervisor*
 - *Reasonably fast* now: hardware (e.g., Intel VT-x) and software-assisted (*Paravirtualization*)
- Advantages
 - **Tight Isolation** (CPU and memory access, virtualized network, virtualized storage, ...)
 - **Better Resource Utilization:** A single host server can manage multiple guest VMs
 - **Better Scalability:** both Vertical (allocate more host resources to guest) and Horizontal (spawn more guests)
 - **Better Disaster Recovery**
 - VM crashed? Just spawn a new one instead
 - VMs have transparent access to replicated network storage
- Drawbacks
 - **You bring *Everything but the kitchen sink*:** Full OS image + all packages + your app (**orders of magnitude** smaller)
 - **Expensive VM setup and teardown** (tens of seconds..minutes, depending on your workload and resources available)
 - **Some performance degradation** (top → Steal time)

Containers! Containers! Containers!



Containers: Concept

- Hypervisors virtualize hardware and then run a full-fledged real OS on it
- But modern OSs **already have powerful primitives** for process, filesystem, network isolation and resource management!

Container includes all relevant (software and data, **NOT** OS!) dependencies for your application/service

...unlike VMs, you cannot *e.g.* run on a fixed specific kernel version.
But this is rarely required for most common apps and services anyway

Containers: Packaging & Delivery

- Linux Containers were intended for **standard packaging and delivery** of software
 - To [re]create development environment
 - To run multiple different versions concurrently
- It is in contrast to e.g. FreeBSD Jails, which were designed a security feature first and foremost
- Docker Images, OCI
 - Optimized for efficient downloads and caching:
 - Layered filesystem
 - Hash-based layer identity
 - Images are tagged (typically with version). latest tag (convention)
- Container Registries (similar to Maven artifact repositories)
 - DockerHub
 - Private installations
 - Cloud offerings from all major cloud providers, incl. Yandex.Cloud

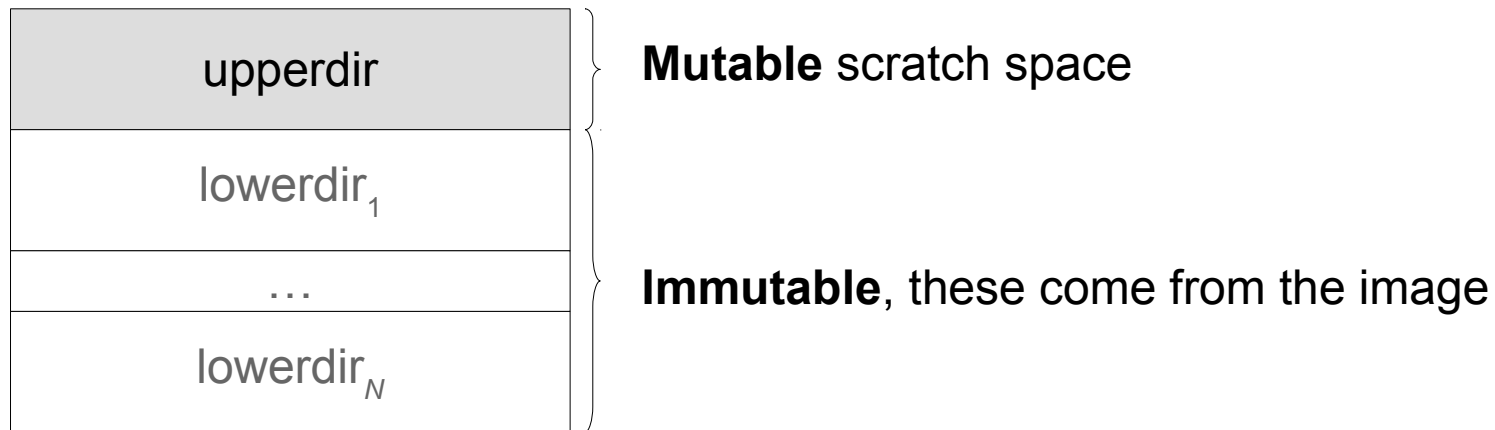
Containers: Isolation → Namespaces

Unit of isolation in Containers is the **Process**

- Namespaces limit **visibility** of sensitive system entities: processes, network interfaces, users, mount points, ...
- **Process isolation**: main process in container runs as PID 1, but its PID in the host is different
- **Network isolation**: container has limited access to host network interfaces, can have special *bridge* network interfaces etc.
- **User isolation**: process inside container runs as root (or as some user that you specify)
- **Mount point isolation**: containers cannot access host storage unless explicitly specified

Containers: Isolation → FS

- `pivot_root` (chroot on steroids) to have / independent of host root filesystem
- Union Filesystem, *overlayfs*



Containers: Limits and Security

- cgroups (Control Groups): Hierarchical Resource Accounting and Limits
 - Limit CPU core usage
 - Memory usage (RSS, resident set size)
 - I/O (read and write iops)
 - Network bandwidth
 - Process is typically killed or throttled when resource overuse is detected
- Container Security: Principle of least privilege
 - Drop capabilities (CAP_...)
 - seccomp-bpf: Selective filtering of syscalls

Container Orchestration

- Containers are much more lightweight than VMs
 - You can have tens and even hundreds of them running on a modest VM!
- Container Orchestration (*E.g.* Kubernetes, Docker Swarm, Nomad, Amazon ECS, ...): Managing lots of containers, and Clusters of similar containers
 - Resource Allocator
 - Workload Scheduler
 - Jobs: batch, throughput-oriented workloads
 - Services (stateless and stateful): interactive, latency-oriented workloads
 - Resilience: restarts/retries, container migration
 - Autoscaling
 - Persistent storage management
- Container Management Philosophy: Cattle **NOT** Pets!

Some Container Solutions

- **Basic Container Management: *Docker*** (safer, simpler production alternative is *rkt*)
 - build image, push (publish to registry), pull (download from registry)
 - run image/run command inside image
 - view stdout/stderr of container process/route process logs to syslog and whatever
- **Basic Orchestration: *Docker Compose***
- Advanced Orchestration: *Hashicorp Nomad, Kubernetes*
- Container-Optimized Linux distributions (*Alpine, distroless*):
Minimal dependencies required to run containers →
→ Smaller attack surface and better performance
- Container Operating Systems: CoreOS, AWS Bottlerocket (experimental)